

Comprehensive PostgreSQL SQL Guide

Table of Contents

1. [SQL Command Types](#)
 - [DDL \(Data Definition Language\)](#)
 - [DML \(Data Manipulation Language\)](#)
 - [DCL \(Data Control Language\)](#)
 - [TCL \(Transaction Control Language\)](#)
 2. [SQL Query Execution Order](#)
 3. [Database Management](#)
 4. [Creating and Querying Tables](#)
 5. [Filtering Data](#)
 6. [Joins](#)
 7. [Grouping Data](#)
 8. [Set Operations](#)
 9. [Modifying Data](#)
-

SQL Command Types

DDL (Data Definition Language) Commands

These commands are used to define and modify the structure of database objects like tables, indexes, and schemas.

- `CREATE DATABASE` : Creates a new database
- `DROP DATABASE` : Deletes a database
- `CREATE TABLE` : Creates a new table
- `DROP TABLE` : Deletes a table
- `ALTER TABLE` : Modifies the structure of an existing table
- `CREATE INDEX` : Creates an index on a table column
- `DROP INDEX` : Deletes an index
- `CREATE SCHEMA` : Creates a new schema
- `DROP SCHEMA` : Deletes a schema

DML (Data Manipulation Language) Commands

These commands are used to manipulate data within the database.

- `INSERT INTO` : Inserts new rows into a table
- `UPDATE` : Updates existing rows in a table
- `DELETE` : Deletes rows from a table
- `SELECT` : Retrieves data from one or more tables
- `TRUNCATE` : Removes all rows from a table without logging individual row deletions

DCL (Data Control Language) Commands

These commands are used to control access to data.

- `GRANT` : Grants privileges to a user or role
- `REVOKE` : Revokes privileges from a user or role

TCL (Transaction Control Language) Commands

These commands are used to manage transactions in the database.

- `BEGIN` : Starts a new transaction
 - `COMMIT` : Commits the current transaction
 - `ROLLBACK` : Rolls back the current transaction to the last commit
-

SQL Query Execution Order

The execution order of SQL clauses in PostgreSQL is different from the written order:

1. **FROM** (and JOINS)
 - Identifies tables involved and performs necessary joins
 - Retrieves raw data from tables
2. **WHERE**
 - Filters rows based on specified conditions

- Only rows satisfying the WHERE condition are included

3. GROUP BY

- Groups filtered rows based on specified columns
- Necessary for aggregate functions (COUNT, SUM, AVG, etc.)

4. HAVING

- Filters groups based on conditions applied to aggregate values
- Only groups satisfying the HAVING condition are included

5. SELECT

- Determines which columns to include in the final result
- Evaluates expressions or calculations specified

6. DISTINCT

- Removes duplicate rows from the result set

7. ORDER BY

- Sorts the result set based on specified columns

8. LIMIT / OFFSET

- Restricts number of rows returned
- Skips specified number of rows

Example Query

```
SELECT department_id, AVG(salary) AS avg_salary
FROM employees
WHERE salary > 40000
GROUP BY department_id
HAVING AVG(salary) > 50000
ORDER BY avg_salary DESC
LIMIT 2;
```

Common Mistakes

1. Using SELECT aliases in WHERE:

```
SELECT employee_name AS name
FROM employees
WHERE name = 'Alice'; -- Error: "name" is not recognized in WHERE
```

Fix: Use the original column name in the WHERE clause.

2. Using HAVING without GROUP BY: HAVING is used to filter groups, so it requires a GROUP BY clause (unless using it with aggregate functions on the entire table).

Database Management

Create a Database

```
CREATE DATABASE database_name;
```

Delete (Drop) a Database

```
DROP DATABASE database_name;
```

Alter a Database

Rename a Database:

```
ALTER DATABASE old_database_name RENAME TO new_database_name;
```

Change the Owner of a Database:

```
ALTER DATABASE old_database_owner OWNER TO new_database_owner;
```

Creating and Querying Tables

Create a New Table

```
CREATE TABLE users(  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(50),  
    age INT,  
    city VARCHAR(50)  
);
```

Insert Data into the Table

```
INSERT INTO users(name, age, city) VALUES  
( 'A', 30, 'Thanjavur'),  
( 'B', 25, 'Kumbakonam'),  
( 'C', 35, 'Orathanadu'),  
( 'D', 45, 'Chennai'),  
( 'E', 30, 'Madurai');
```

Retrieve Data Using SQL Commands

SELECT Statement

```
SELECT * FROM users;  
  
SELECT id, name, age, city FROM users;
```

Column Aliases

```
SELECT name AS user_name, age AS user_age FROM users;
```

ORDER BY Clause

```
SELECT * FROM users ORDER BY age ASC;  
SELECT * FROM users ORDER BY age DESC;
```

SELECT DISTINCT

```
SELECT DISTINCT age FROM users;           -- Works  
SELECT DISTINCT age, name FROM users;     -- Works for unique combinations
```

Filtering Data

Sample Table Setup

```
CREATE TABLE employees(  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(100),  
    department VARCHAR(50),  
    salary NUMERIC(10,2),  
    hire_date DATE  
);  
  
INSERT INTO employees(name, department, salary, hire_date) VALUES  
( 'Alice', 'HR', 50000, '2020-06-15'),  
( 'Bob', 'Engineering', 75000, '2019-03-22'),  
( 'Charlie', 'Marketing', 60000, '2021-08-10'),  
( 'David', 'Engineering', 80000, '2018-11-05'),  
( 'Eve', 'HR', 55000, '2022-01-14'),  
( 'Frank', 'Marketing', NULL, '2020-07-01');
```

Filtering Examples

WHERE Clause

Filter rows based on a specified condition:

```
SELECT * FROM employees WHERE department = 'Engineering';
```

AND Operator

Combines two boolean expressions and returns true if both are true:

```
SELECT * FROM employees WHERE department = 'Engineering' AND salary >= 75000;
```

OR Operator

Returns true if either expression is true:

```
SELECT * FROM employees WHERE department = 'HR' OR department = 'Marketing';
```

LIMIT Clause

Retrieves a subset of rows:

```
SELECT * FROM employees LIMIT 3;  
SELECT * FROM employees ORDER BY id DESC LIMIT 3;
```

FETCH Clause

Similar to LIMIT:

```
SELECT * FROM employees FETCH FIRST 2 ROWS ONLY;
```

IN Operator

Selects data matching any value in a list:

```
SELECT * FROM employees WHERE department IN ('HR', 'Marketing');
```

BETWEEN Operator

Selects data within a range:

```
SELECT * FROM employees WHERE salary BETWEEN 50000 AND 70000;
```

LIKE Operator

Pattern matching for strings:

```
SELECT * FROM employees WHERE name LIKE 'A%';      -- Names starting with A
SELECT * FROM employees WHERE name LIKE '%e';      -- Names ending with e
SELECT * FROM employees WHERE name LIKE 'A_____'; -- Names starting with A and exactly 5 characters long
SELECT * FROM employees WHERE name ILIKE 'a%';      -- Case-insensitive match for names starting with a or A
SELECT * FROM employees WHERE name LIKE '%e%';      -- Names containing e
SELECT * FROM employees WHERE name NOT LIKE 'A%';   -- Names not starting with A
```

IS NULL Operator

Checks for NULL values:

```
SELECT * FROM employees WHERE salary IS NULL;
SELECT * FROM employees WHERE salary IS NOT NULL;
```

OFFSET

Skips specified number of rows:

```
SELECT * FROM employees OFFSET 2 ROWS;
SELECT * FROM employees OFFSET 2 ROWS LIMIT 2;
SELECT * FROM employees OFFSET 2 ROWS FETCH FIRST 2 ROWS ONLY;
```

Joins

Sample Tables Setup

```
DROP TABLE IF EXISTS employees;

CREATE TABLE employees (
    employee_id SERIAL PRIMARY KEY,
    employee_name VARCHAR(100) NOT NULL,
    department_id INT,
    salary NUMERIC(10, 2)
);

CREATE TABLE departments(
    department_id SERIAL PRIMARY KEY,
    department_name VARCHAR(50) NOT NULL
);

INSERT INTO employees (employee_name, department_id, salary) VALUES
('Alice', 1, 50000),
('Bob', 2, 60000),
('Charlie', 1, 55000),
('David', NULL, 70000),
('Eve', 3, 45000);

INSERT INTO departments (department_name) VALUES
('HR'),
('Engineering'),
('Marketing');
```

Table Aliases

```
SELECT e.employee_name, d.department_name
FROM employees AS e
JOIN departments AS d
ON e.department_id = d.department_id;
```

INNER JOIN

Returns only rows with matching values in both tables:

```
SELECT e.employee_name, d.department_name
FROM employees AS e
INNER JOIN departments AS d
ON e.department_id = d.department_id;
```

LEFT JOIN

Returns all rows from the left table and matching rows from the right table:

```
SELECT e.employee_name, d.department_name
FROM employees AS e
LEFT JOIN departments AS d
ON e.department_id = d.department_id;
```

Example with tables reversed:

```
SELECT e.employee_name, d.department_name
FROM departments AS d
LEFT JOIN employees AS e
ON d.department_id = e.department_id;
```

RIGHT JOIN

Returns all rows from the right table and matching rows from the left table:

```
SELECT e.employee_name, d.department_name
FROM employees AS e
RIGHT JOIN departments AS d
ON e.department_id = d.department_id;
```

Example with tables reversed:

```
SELECT e.employee_name, d.department_name
FROM departments AS d
RIGHT JOIN employees AS e
ON d.department_id = e.department_id;
```

Grouping Data

GROUP BY Clause

Groups rows that have the same values in specified columns into summary rows.

Syntax:

```
SELECT column1, aggregate_function(column2)
FROM table_name
GROUP BY column1;
```

Example 1: Basic GROUP BY

Find the total salary for each department:

```
SELECT department_id, SUM(salary)
FROM employees
GROUP BY department_id;
```

Example 2: GROUP BY with JOIN

Find the total salary for each department with department name:

```
SELECT e.department_id, d.department_name, SUM(salary) AS TotalSalary
FROM employees AS e
INNER JOIN departments AS d ON e.department_id = d.department_id
GROUP BY e.department_id, d.department_name;
```

Example 3: COUNT

Count the number of employees in each department:

```
SELECT e.department_id, d.department_name, COUNT(salary) AS TotalEmployee
FROM employees AS e
INNER JOIN departments AS d ON e.department_id = d.department_id
GROUP BY e.department_id, d.department_name;
```

Example 4: AVG

Calculate the average salary for each department:

```
SELECT e.department_id, d.department_name, AVG(salary) as department_AvgSalary
FROM employees AS e
INNER JOIN departments AS d ON e.department_id = d.department_id
GROUP BY e.department_id, d.department_name;
```

Example 5: MIN

Find the minimum salary in each department:

```
SELECT e.department_id, d.department_name, MIN(salary) as department_MinSalary
FROM employees AS e
INNER JOIN departments AS d ON e.department_id = d.department_id
GROUP BY e.department_id, d.department_name;
```

Example 6: MAX

Find the maximum salary in each department:

```
SELECT e.department_id, d.department_name, MAX(salary) as department_MaxSalary
FROM employees AS e
INNER JOIN departments AS d ON e.department_id = d.department_id
GROUP BY e.department_id, d.department_name;
```

HAVING Clause

Filters groups based on a condition applied to aggregate function results.

Syntax:

```
SELECT column1, aggregate_function(column2)
FROM table_name
GROUP BY column1
HAVING condition;
```

Example 1: HAVING with COUNT

Find departments with more than 1 employee:

```
SELECT d.department_id, department_name, COUNT(salary) AS total_employees
FROM employees AS e
JOIN departments AS d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING COUNT(e.department_id) >= 2;
```

Example 2: HAVING with SUM

Find departments where the total salary is greater than 100,000:

```
SELECT d.department_id, d.department_name, SUM(salary) AS TotalSalary
FROM employees AS e
JOIN departments AS d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING SUM(salary) >= 100000;
```

Example 3: HAVING with AVG

Find departments where the average salary is greater than 10,000:

```
SELECT d.department_id, d.department_name, AVG(salary) AS average_salary
FROM employees AS e
JOIN departments AS d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING AVG(salary) >= 10000;
```

Example 4: HAVING with MIN

Find departments where the minimum salary is less than 50,000:

```
SELECT d.department_id, d.department_name, MIN(salary) AS min_salary
FROM employees AS e
JOIN departments AS d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING MIN(salary) < 50000;
```

Example 5: HAVING with MAX

Find departments where the maximum salary is greater than 50,000:

```
SELECT d.department_id, d.department_name, MAX(salary) AS max_salary
FROM employees AS e
JOIN departments AS d ON e.department_id = d.department_id
GROUP BY d.department_id, d.department_name
HAVING MAX(salary) > 50000;
```

Set Operations

Sample Tables Setup


```
CREATE TABLE students (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(100) NOT NULL  
);  
  
CREATE TABLE teachers (  
    id SERIAL PRIMARY KEY,  
    name VARCHAR(100) NOT NULL  
);  
  
INSERT INTO teachers (name) VALUES  
( 'Alice'), ( 'Bob'), ( 'Tamil'), ( 'Thillai');  
  
INSERT INTO students (name) VALUES  
( 'Alice'), ( 'Bob'), ( 'sabi'), ( 'sathish');
```

UNION

Combines results of two or more SELECT queries into a single result set, removing duplicates:

```
SELECT * FROM teachers  
UNION  
SELECT * FROM students;
```

With duplicates:

```
SELECT * FROM teachers  
UNION ALL  
SELECT * FROM students;
```

INTERSECT

Returns common rows that appear in both SELECT queries:

```
SELECT * FROM teachers  
INTERSECT  
SELECT * FROM students;
```

EXCEPT

Returns rows from the first SELECT query that are not present in the second:

```
SELECT * FROM teachers  
EXCEPT  
SELECT * FROM students;
```

Summary

- **UNION:** Combines results from multiple queries and removes duplicates
- **INTERSECT:** Returns common rows from multiple queries
- **EXCEPT:** Returns rows from the first query that are not in the second query

Modifying Data

Create a Table

```
CREATE TABLE worker
(
    id SERIAL PRIMARY KEY,
    name VARCHAR(50),
    position VARCHAR(50)
);
```

INSERT

Insert a single row:

```
INSERT INTO worker (name, position) VALUES ('Abi', 'Software Engineer');
```

Insert multiple rows:

```
INSERT INTO worker (name, position) VALUES
('Dharshini', 'Data Scientist'),
('Thilai', 'Product Manager'),
('Tamizh', 'CEO');
```

UPDATE

Update existing rows:

```
UPDATE worker SET position = 'Senior Software Engineer' WHERE id = 1;
```

DELETE

Delete rows:

```
DELETE FROM worker WHERE id = 3;
```

UPSERT

Insert or update if conflict:

```
INSERT INTO worker (name, position)
VALUES ('Thillai', 'Team Leader')
ON CONFLICT(name)
DO UPDATE SET position = EXCLUDED.position;
```

Note: The ON CONFLICT clause requires a unique identifier to detect conflicts. If you try to use ON CONFLICT without a primary key or unique constraint, PostgreSQL will throw an error.